

Hierarchical Clustering

Osama Bin Zahir

2023-11-30

#Loading the dataset

```
data=read.csv("C:\\Users\\Osama Zahir\\Downloads\\Cereals.csv")
head(data)
```

```
##              name mfr type calories protein fat sodium fiber carbo
## 1          100%_Bran  N   C        70         4   1   130  10.0   5.0
## 2      100%_Natural_Bran  Q   C       120         3   5    15   2.0   8.0
## 3           All-Bran   K   C        70         4   1   260   9.0   7.0
## 4 All-Bran_with_Extra_Fiber  K   C        50         4   0   140  14.0   8.0
## 5          Almond_Delight  R   C       110         2   2   200   1.0  14.0
## 6 Apple_Cinnamon_Cheerios  G   C       110         2   2   180   1.5  10.5
##  sugars potass vitamins shelf weight cups   rating
## 1      6    280      25     3      1 0.33 68.40297
## 2      8    135       0     3      1 1.00 33.98368
## 3      5    320      25     3      1 0.33 59.42551
## 4      0    330      25     3      1 0.50 93.70491
## 5      8     NA      25     3      1 0.75 34.38484
## 6     10     70      25     1      1 0.75 29.50954
```

#Viewing the summary and structure of that dataset

```
summary(data)
```

```
##      name      mfr      type      calories
## Length:77      Length:77      Length:77      Min.   : 50.0
## Class :character Class :character Class :character 1st Qu.:100.0
## Mode  :character Mode  :character Mode  :character Median :110.0
##                                     Mean   :106.9
##                                     3rd Qu.:110.0
##                                     Max.   :160.0
##
##      protein      fat      sodium      fiber
## Min.   :1.000      Min.   :0.000      Min.   : 0.0      Min.   : 0.000
## 1st Qu.:2.000      1st Qu.:0.000      1st Qu.:130.0      1st Qu.: 1.000
## Median :3.000      Median :1.000      Median :180.0      Median : 2.000
## Mean   :2.545      Mean   :1.013      Mean   :159.7      Mean   : 2.152
## 3rd Qu.:3.000      3rd Qu.:2.000      3rd Qu.:210.0      3rd Qu.: 3.000
## Max.   :6.000      Max.   :5.000      Max.   :320.0      Max.   :14.000
##
##      carbo      sugars      potass      vitamins
## Min.   : 5.0      Min.   : 0.000      Min.   : 15.00      Min.   : 0.00
## 1st Qu.:12.0      1st Qu.: 3.000      1st Qu.: 42.50      1st Qu.: 25.00
## Median :14.5      Median : 7.000      Median : 90.00      Median : 25.00
## Mean   :14.8      Mean   : 7.026      Mean   : 98.67      Mean   : 28.25
## 3rd Qu.:17.0      3rd Qu.:11.000      3rd Qu.:120.00      3rd Qu.: 25.00
## Max.   :23.0      Max.   :15.000      Max.   :330.00      Max.   :100.00
## NA's   :1        NA's   :1        NA's   :2
##      shelf      weight      cups      rating
## Min.   :1.000      Min.   :0.50      Min.   :0.250      Min.   :18.04
## 1st Qu.:1.000      1st Qu.:1.00      1st Qu.:0.670      1st Qu.:33.17
## Median :2.000      Median :1.00      Median :0.750      Median :40.40
## Mean   :2.208      Mean   :1.03      Mean   :0.821      Mean   :42.67
## 3rd Qu.:3.000      3rd Qu.:1.00      3rd Qu.:1.000      3rd Qu.:50.83
## Max.   :3.000      Max.   :1.50      Max.   :1.500      Max.   :93.70
##
```

```
str(data)
```

```
## 'data.frame': 77 obs. of 16 variables:
## $ name : chr "100%_Bran" "100%_Natural_Bran" "All-Bran" "All-Bran_with_Extra_Fiber" ...
## $ mfr : chr "N" "Q" "K" "K" ...
## $ type : chr "C" "C" "C" "C" ...
## $ calories: int 70 120 70 50 110 110 110 130 90 90 ...
## $ protein : int 4 3 4 4 2 2 2 3 2 3 ...
## $ fat : int 1 5 1 0 2 2 0 2 1 0 ...
## $ sodium : int 130 15 260 140 200 180 125 210 200 210 ...
## $ fiber : num 10 2 9 14 1 1.5 1 2 4 5 ...
## $ carbo : num 5 8 7 8 14 10.5 11 18 15 13 ...
## $ sugars : int 6 8 5 0 8 10 14 8 6 5 ...
## $ potass : int 280 135 320 330 NA 70 30 100 125 190 ...
## $ vitamins: int 25 0 25 25 25 25 25 25 25 25 ...
## $ shelf : int 3 3 3 3 3 1 2 3 1 3 ...
## $ weight : num 1 1 1 1 1 1 1 1.33 1 1 ...
## $ cups : num 0.33 1 0.33 0.5 0.75 0.75 1 0.75 0.67 0.67 ...
## $ rating : num 68.4 34 59.4 93.7 34.4 ...
```

#Loading required packages

```
library(caret)
```

```
## Loading required package: ggplot2
```

```
## Loading required package: lattice
```

```
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 4.3.2
```

```
## corrplot 0.92 loaded
```

```
library(ggcorrplot)
```

```
## Warning: package 'ggcorrplot' was built under R version 4.3.2
```

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.3.2
```

```
## Warning: package 'forcats' was built under R version 4.3.2
```

```
## — Attaching core tidyverse packages — tidyverse 2.0.0 —
## ✓ dplyr      1.1.3    ✓ readr      2.1.4
## ✓ forcats    1.0.0    ✓ stringr    1.5.0
## ✓ lubridate  1.9.2    ✓ tibble     3.2.1
## ✓ purrr      1.0.2    ✓ tidyr      1.3.0
```

```
## — Conflicts — tidyverse_conflicts() —
## ✗ dplyr::filter() masks stats::filter()
## ✗ dplyr::lag()     masks stats::lag()
## ✗ purrr::lift()    masks caret::lift()
## ⓘ Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
library(tidyr)
library(dplyr)
library(e1071)
library(reshape2)
```

```
##
## Attaching package: 'reshape2'
##
## The following object is masked from 'package:tidyr':
##
##   smiths
```

```
library(factoextra)
```

```
## Warning: package 'factoextra' was built under R version 4.3.2
```

```
## Welcome! Want to learn more? See two factoextra-related books at https://goo.gl/ve3WBa
```

```
library(cluster)
library(cowplot)
```

```
## Warning: package 'cowplot' was built under R version 4.3.2
```

```
##
## Attaching package: 'cowplot'
##
## The following object is masked from 'package:lubridate':
##
##   stamp
```

```
library(pander)
library(kernlab)
```

```
##  
## Attaching package: 'kernlab'  
##  
## The following object is masked from 'package:purrr':  
##  
##     cross  
##  
## The following object is masked from 'package:ggplot2':  
##  
##     alpha
```

```
library(FactoMineR)
```

```
## Warning: package 'FactoMineR' was built under R version 4.3.2
```

#Data Preprocessing. Remove all cereals with missing values.

```
dim(data)
```

```
## [1] 77 16
```

```
c_d2=na.omit(data)  
dim(c_d2)
```

```
## [1] 74 16
```

#There were 4 missing values in the dataset

#Assigning row names to the cereal column

```
c_d3 = as.data.frame(c_d2)  
row.names(c_d3) = c_d3[,1]  
c_d4 = c_d3[,-1]
```

#Only selecting numerical values and removing categorical variables

```
c_d5 = c_d4[, c(3:11,13:15)]
```

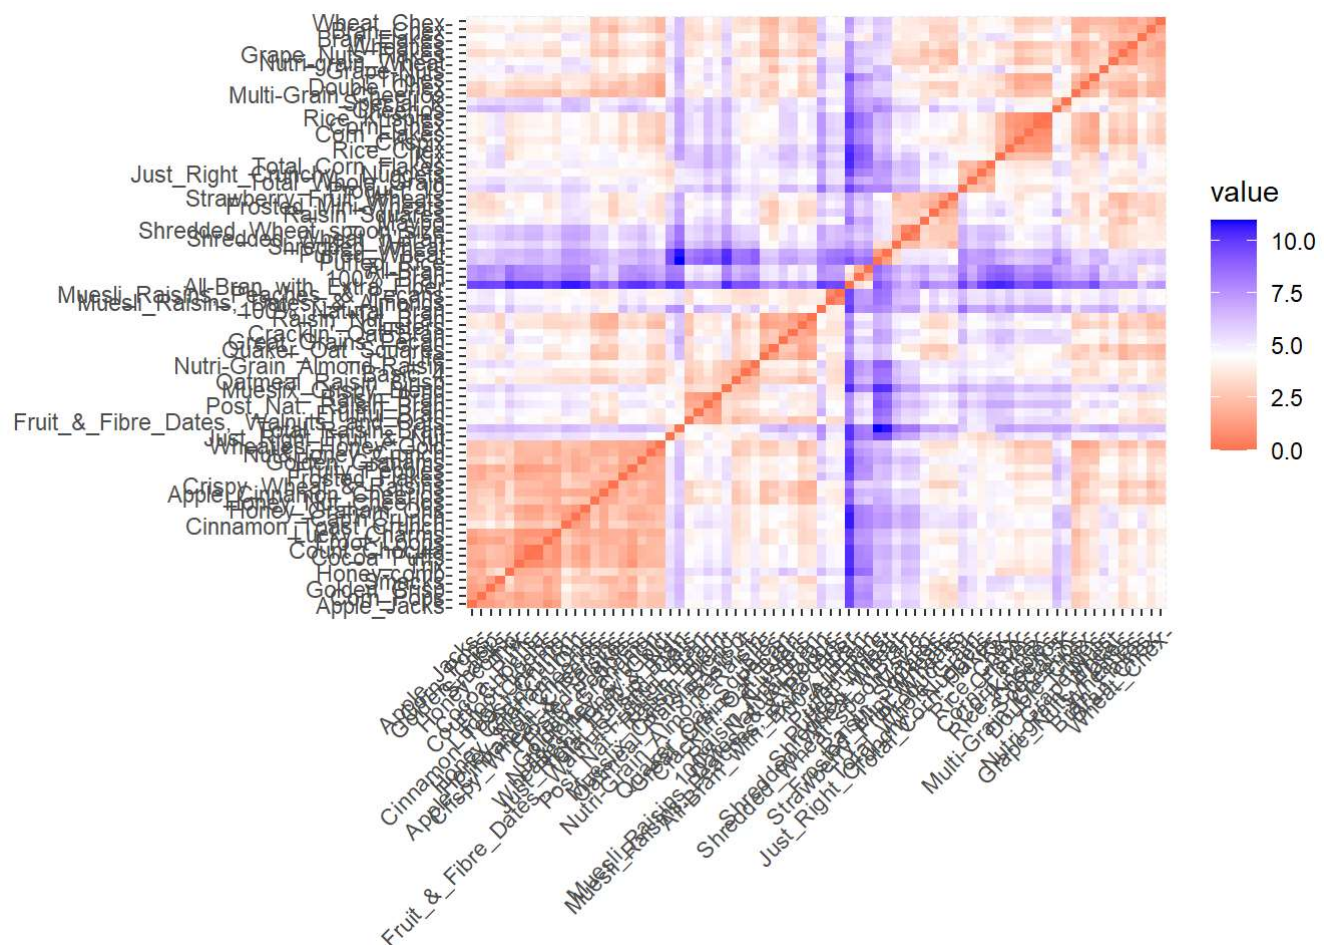
#Normalizing the data using the scale function

```
c_d5 = scale(c_d5)  
head(c_d5)
```

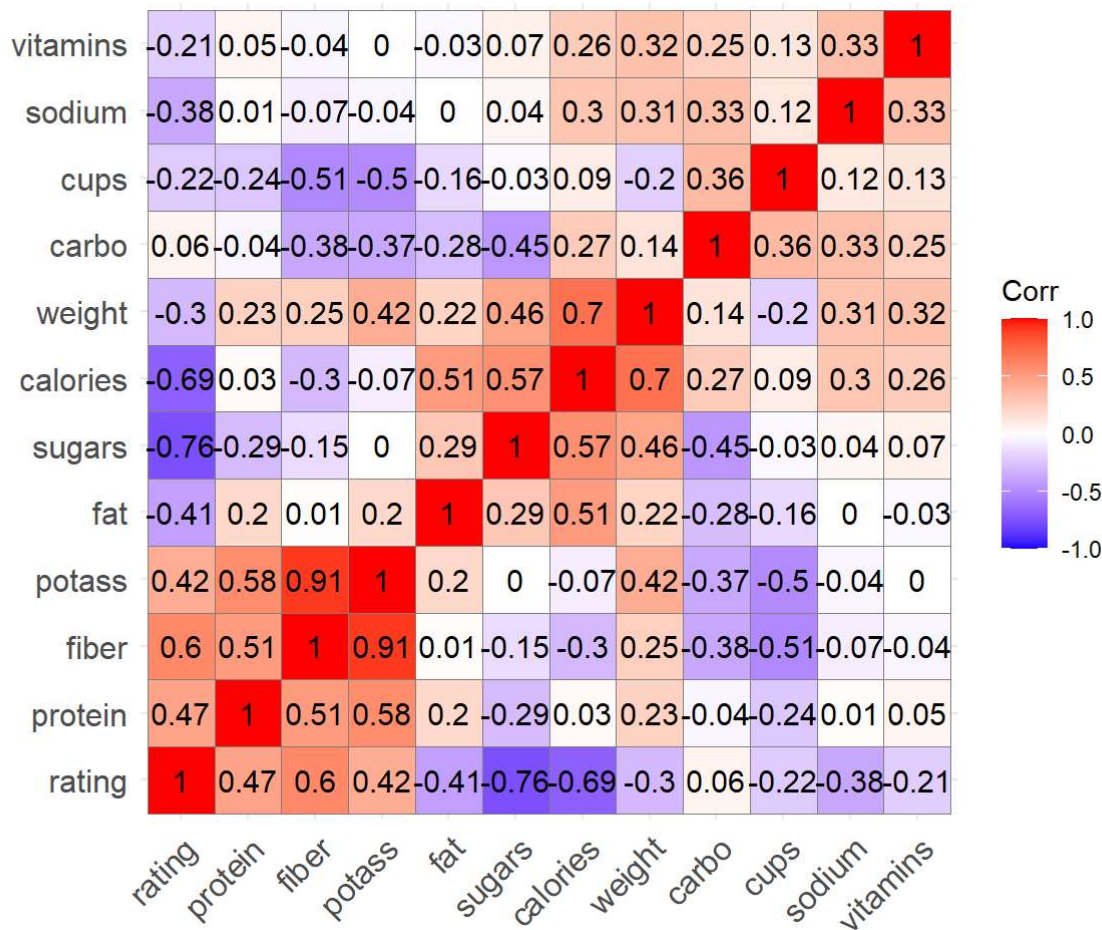
##	calories	protein	fat	sodium
## 100%_Bran	-1.8659155	1.3817478	0.0000000	-0.3910227
## 100%_Natural_Bran	0.6537514	0.4522084	3.9728810	-1.7804186
## All-Bran	-1.8659155	1.3817478	0.0000000	1.1795987
## All-Bran_with_Extra_Fiber	-2.8737823	1.3817478	-0.9932203	-0.2702057
## Apple_Cinnamon_Cheerios	0.1498180	-0.4773310	0.9932203	0.2130625
## Apple_Jacks	0.1498180	-0.4773310	-0.9932203	-0.4514312
##	fiber	carbo	sugars	potass
## 100%_Bran	3.22866747	-2.5001396	-0.2542051	2.5605229
## 100%_Natural_Bran	-0.07249167	-1.7292632	0.2046041	0.5147738
## All-Bran	2.81602258	-1.9862220	-0.4836096	3.1248675
## All-Bran_with_Extra_Fiber	4.87924705	-1.7292632	-1.6306324	3.2659536
## Apple_Cinnamon_Cheerios	-0.27881412	-1.0868662	0.6634132	-0.4022862
## Apple_Jacks	-0.48513656	-0.9583868	1.5810314	-0.9666308
##	vitamins	weight	cups	rating
## 100%_Bran	-0.1818422	-0.2008324	-2.0856582	1.8549038
## 100%_Natural_Bran	-1.3032024	-0.2008324	0.7567534	-0.5977113
## All-Bran	-0.1818422	-0.2008324	-2.0856582	1.2151965
## All-Bran_with_Extra_Fiber	-0.1818422	-0.2008324	-1.3644493	3.6578436
## Apple_Cinnamon_Cheerios	-0.1818422	-0.2008324	-0.3038480	-0.9165248
## Apple_Jacks	-0.1818422	-0.2008324	0.7567534	-0.6553998

Question 1 (part A): Apply hierarchical clustering to the data using Euclidean distance to the normalized measurements and looking at the correaltion values by plotting the corrplot

```
distance_table <- get_dist(c_d5)
fviz_dist(distance_table)
```



```
corr_plot = cor(c_d5)
ggcorrplot(corr_plot, outline.color = "grey50", lab = TRUE, hc.order = TRUE, type = "full")
```

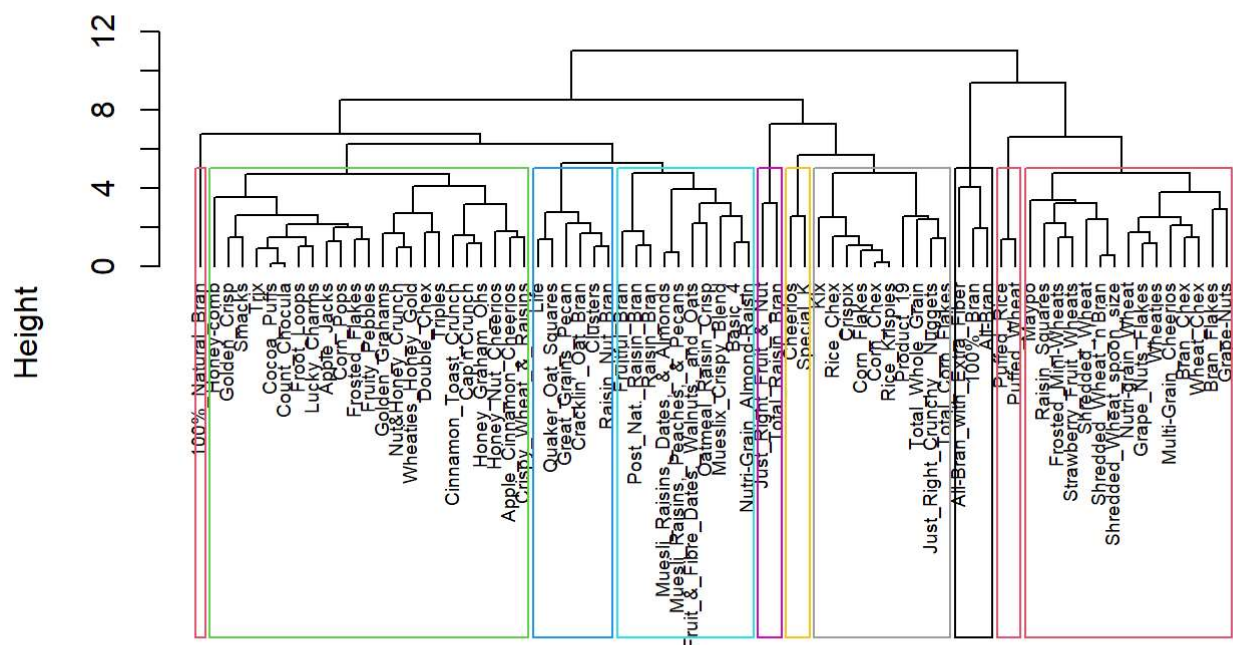


#Sugar and calories are highly negatively correlated with rating. Also, Potass is highly positively correlated with fiber and Protien.

Question 1 (part B): comparing hierarchical clustering with different linkages: single, average, complete and ward.

```
# Hierarchical clustering using Complete Linkage
hc1 <- hclust(distance_table, method = "complete" )
# Plot the obtained dendrogram
plot(hc1, cex = 0.6, hang = -1, main = "Dendrogram of Hierarchical Clustering")
rect.hclust(hc1, k = 10, border = 2:10)
```


Dendrogram of Hierarchical Clustering



distance_table
hclust (*, "complete")

#Computing with AGNES and with different linkage methods

```
hc_single <- agnes(distance_table, method = "single")
print(hc_single$ac)
```

```
## [1] 0.6072384
```

```
hc_complete <- agnes(distance_table, method = "complete")
print(hc_complete$ac)
```

```
## [1] 0.8469328
```

```
hc_average <- agnes(distance_table, method = "average")
print(hc_average$ac)
```

```
## [1] 0.7881955
```

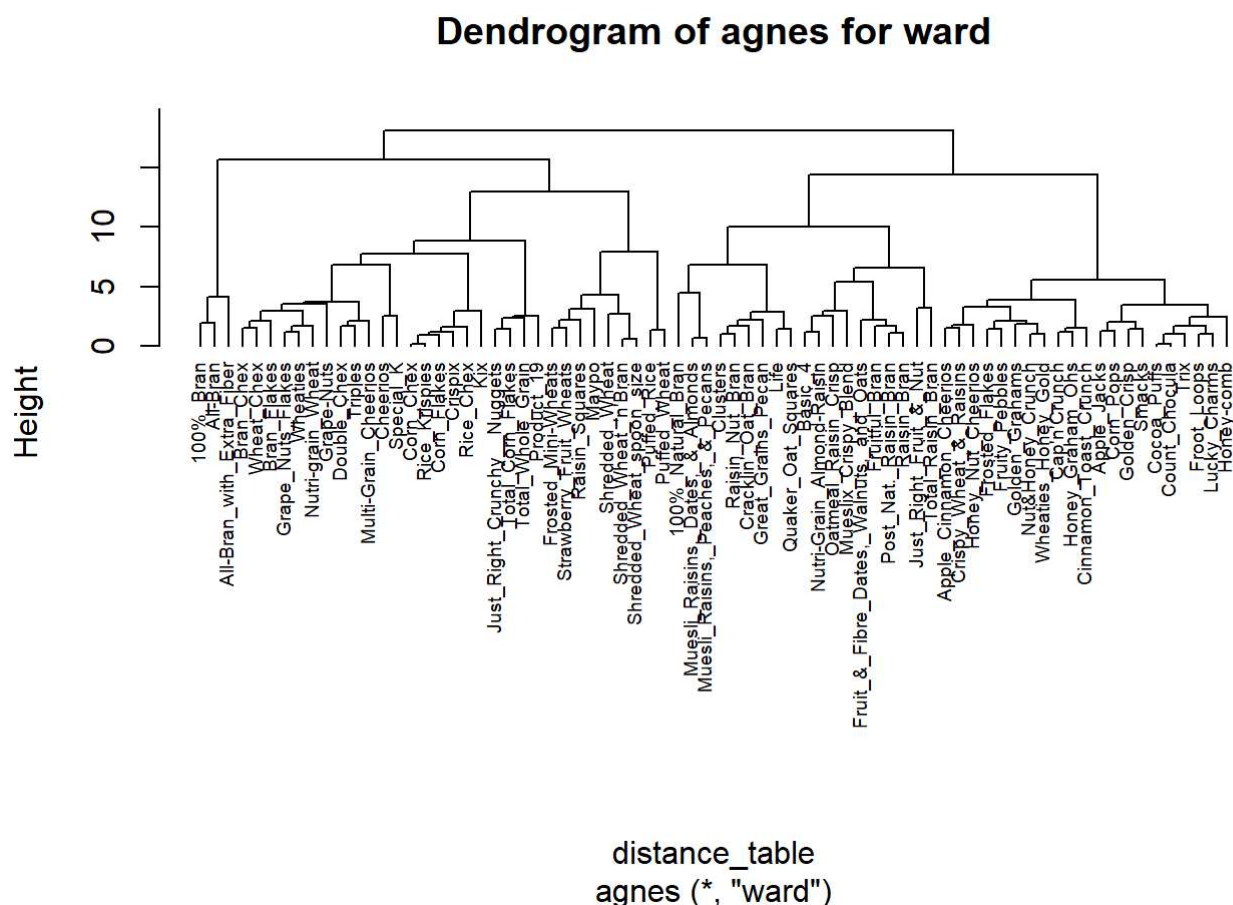
```
hc_ward <- agnes(distance_table, method = "ward")
print(hc_ward$ac)
```

```
## [1] 0.9087265
```

#These results confirm that the Ward linkage, which provides 90.87% accuracy, is the optimal agglomerative (AGNES) linkage to use.

#Visualizing the dendrogram

```
hc_Ward <- agnes(distance_table, method = "ward")
pltree(hc_Ward, cex = 0.6, hang = -1, main = "Dendrogram of agnes for ward")
```



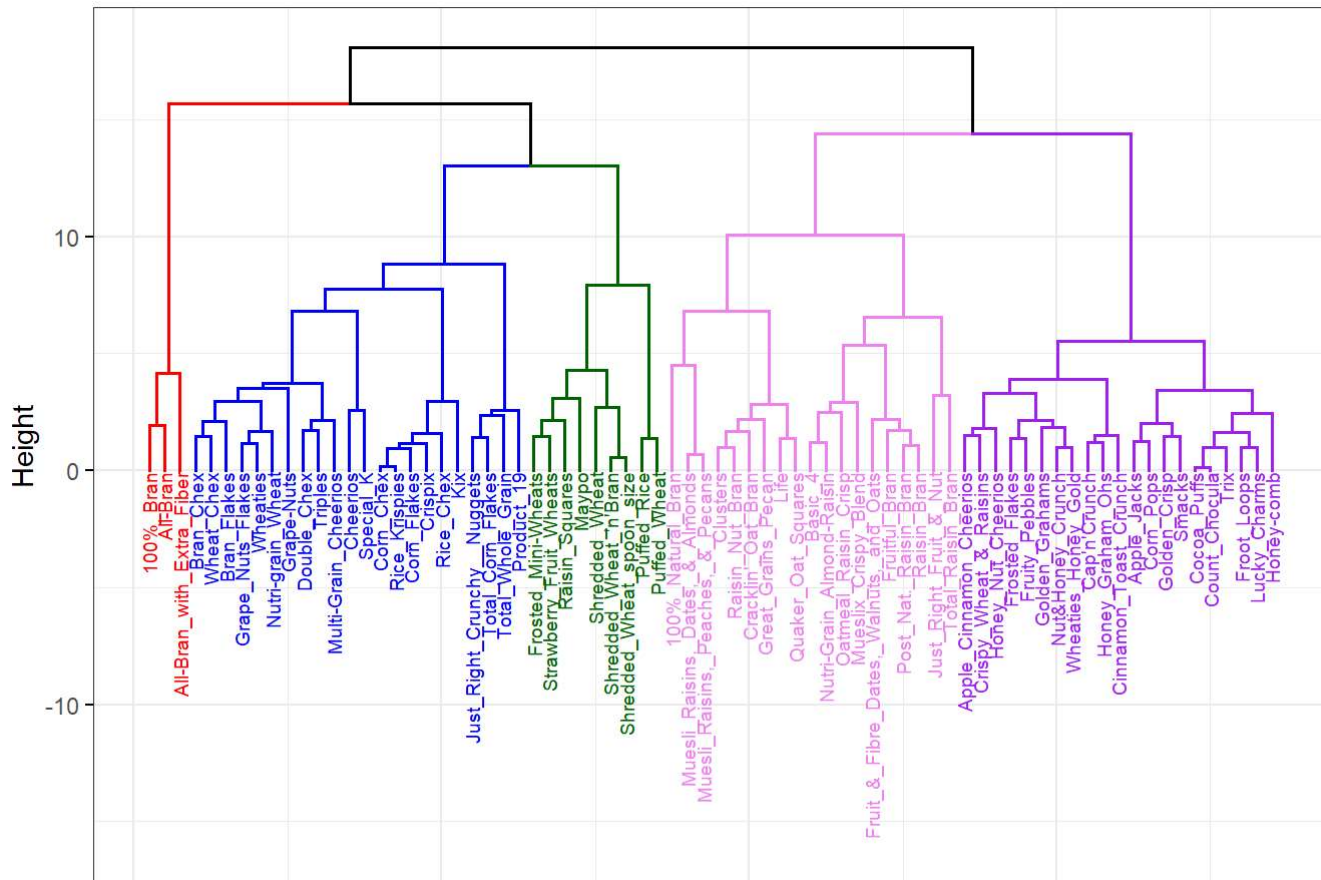
Question 2: How many cluster would you choose?

The largest difference in height can be used to determine the k value hence K =5 is the best option.

```
fviz_dend(hc_ward, k = 5, main = "Dendrogram of AGNES (Ward)", cex = 0.5, k_colors = c("red", "blue", "darkgreen", "violet", "purple"), color_labels_by_k = TRUE, labels_track_height = 16, ggtheme = theme_bw())
```

```
## Warning: The `scale` argument of `guides()` cannot be `FALSE`. Use "none" instead as
## of ggplot2 3.3.4.
## i The deprecated feature was likely used in the factoextra package.
## Please report the issue at <https://github.com/kassambara/factoextra/issues>.
## This warning is displayed once every 8 hours.
## Call `lifecycle::last_lifecycle_warnings()` to see where this warning was
## generated.
```

Dendrogram of AGNES (Ward)



```
c_d6 <- cutree(hc_ward, k = 5)
Clustered df <- as.data.frame(cbind ( c_d5, c_d6 ))
```

Question 3: Comment on the structure of the clusters and on their stability. Hint: To check stability, partition the data and see how well clusters formed based on one part apply to the other part

#For the stability of the clusters, We will partition the data into A and B.

```
cereal_a = c_d5[1:55,]  
cereal_b = c_d5[56:74,]
```

#Computing the distances of cereal_a

```
distance_cereal_a = get_dist(cereal_a)
```

#Compute with AGNES and with different linkage methods for cereal_a

```
hc_single_cereal_a <- agnes(distance_cereal_a, method = "single")  
print(hc_single_cereal_a$ac)
```

```
## [1] 0.6663587
```

```
hc_complete_cereal_a <- agnes(distance_cereal_a, method = "complete")  
print(hc_complete_cereal_a$ac)
```

```
## [1] 0.8285192
```

```
hc_average_cereal_a <- agnes(distance_cereal_a, method = "average")  
print(hc_average_cereal_a$ac)
```

```
## [1] 0.7646836
```

```
hc_ward_cereal_a <- agnes(distance_cereal_a, method = "ward")  
print(hc_ward_cereal_a$ac)
```

```
## [1] 0.8891086
```

#With 88.91% accuracy, it enables us to establish that the best linkage for cereal_a is Ward.

#Computing the distances of cereal_b

```
distance_cereal_b = get_dist(cereal_b)
```

#Compute with AGNES and with different linkage methods for cereal_b

```
hc_single_cereal_b <- agnes(distance_cereal_b, method = "single")  
print(hc_single_cereal_b$ac)
```

```
## [1] 0.4805129
```

```
hc_complete_cereal_b <- agnes(distance_cereal_b, method = "complete")  
print(hc_complete_cereal_b$ac)
```

```
## [1] 0.71298
```

```
hc_average_cereal_b <- agnes(distance_cereal_b, method = "average")
print(hc_average_cereal_b$ac)
```

```
## [1] 0.6232053
```

```
hc_ward_cereal_b <- agnes(distance_cereal_b, method = "ward")
print(hc_ward_cereal_b$ac)
```

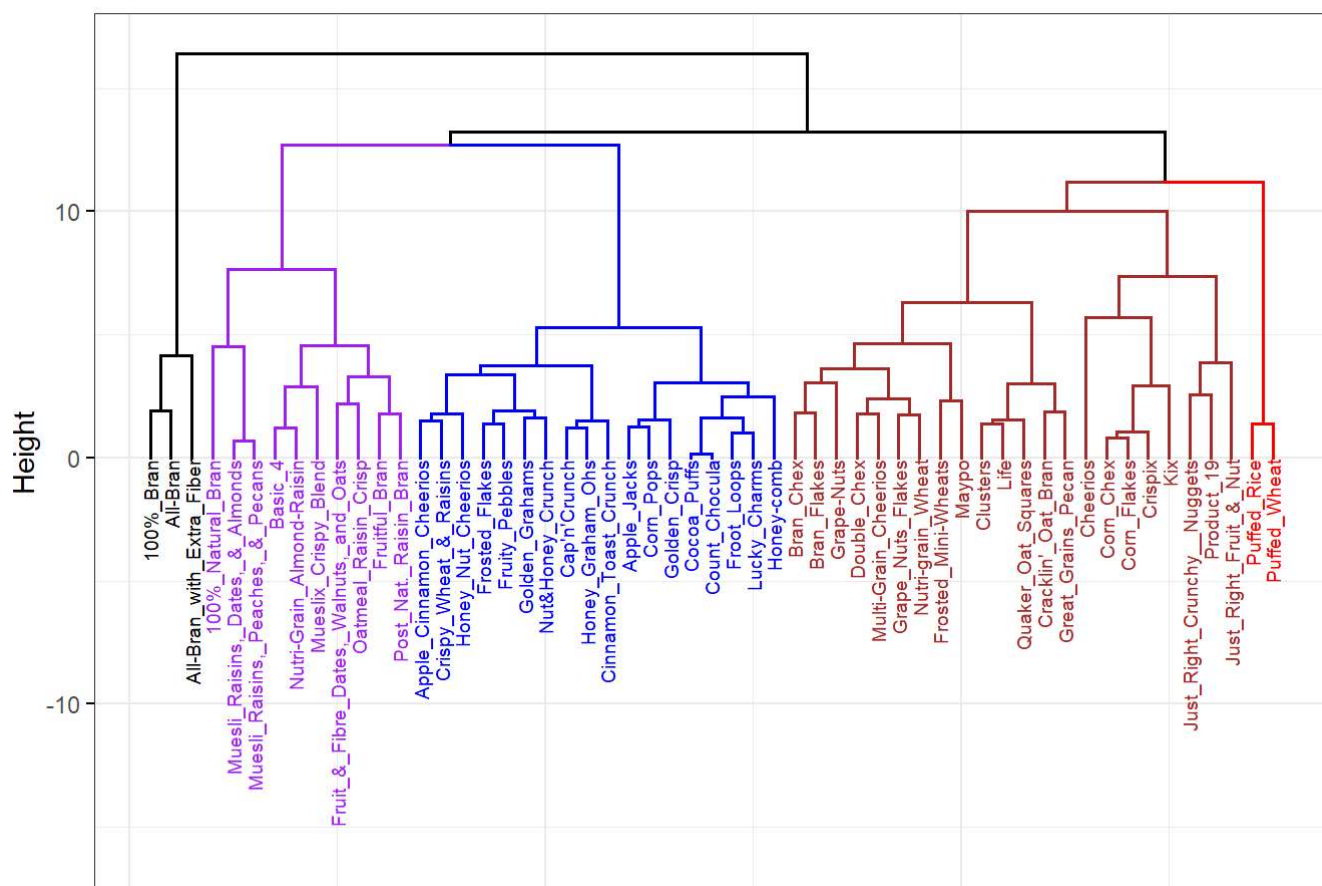
```
## [1] 0.7710122
```

#With 77.10% accuracy, it enables us to establish that the best linkage for cereal_a is Ward.

#Plotting dendrogram of cereal_a and cereal_b

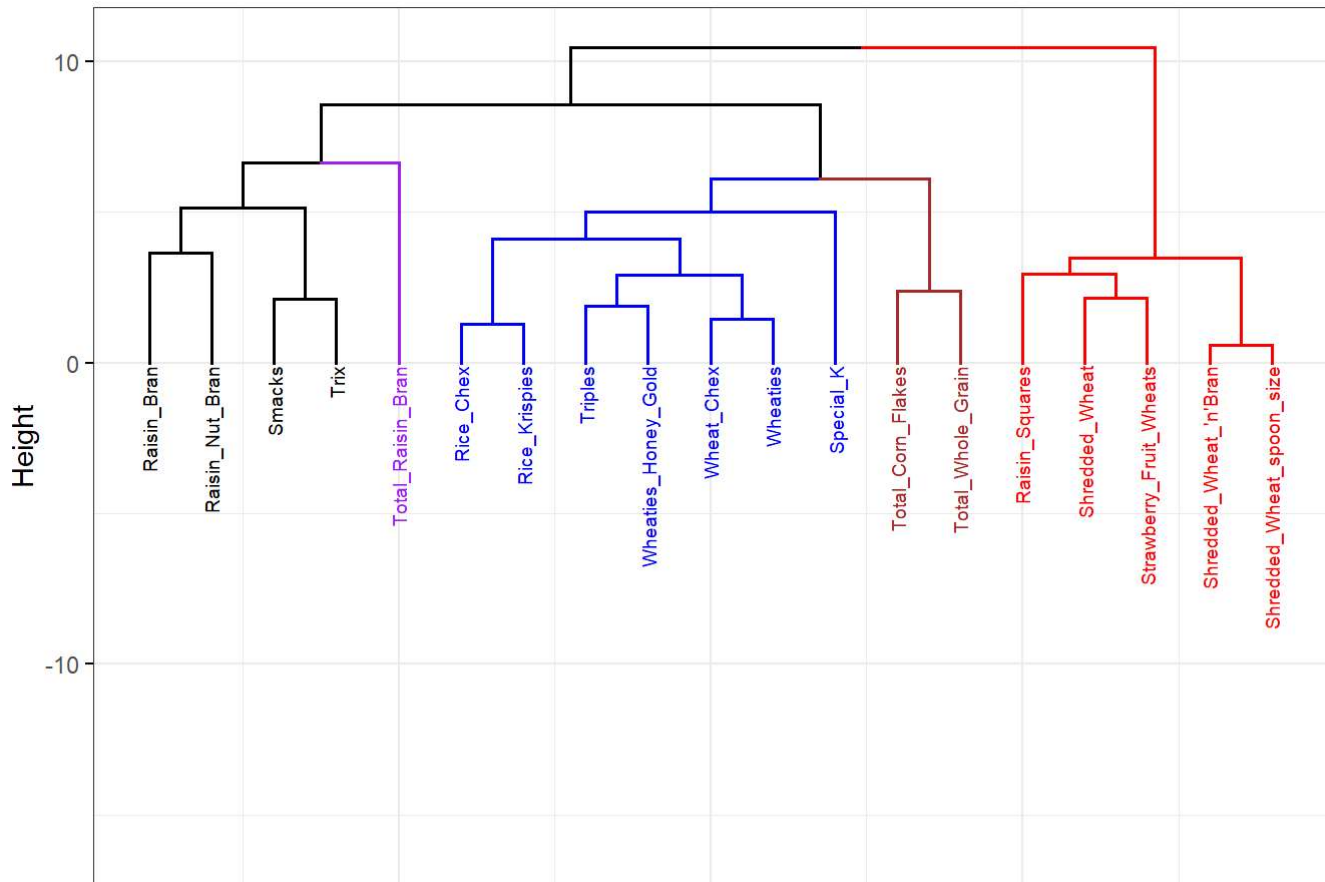
```
fviz_dend(hc_ward_cereal_a, k = 5, main = "Cereal_a Dendrogram of AGNES", cex = 0.5, k_colors = c(
  "black", "purple", "blue", "brown", "red"), color_labels_by_k = TRUE, labels_track_height = 16, g
  theme = theme_bw())
```

Cereal_a Dendrogram of AGNES



```
fviz_dend(hc_ward_cereal_b, k = 5, main = "Cereal_b Dendrogram of AGNES", cex = 0.5, k_colors = c(
  "black", "purple", "blue", "brown", "red"), color_labels_by_k = TRUE, labels_track_height = 16, g
  theme = theme_bw())
```

Cereal_b Dendrogram of AGNES



Question 3 (part B): Use the cluster centroids from A to assign each record in partition B (each record is assigned to the cluster with the closest centroid)

```
Clustered_df_A <- cutree (hc_ward_cereal_a, k=5)
Clusters_A <- as.data.frame(cbind(cereal_a, Clustered_df_A))
Clust_1 <- colMeans (Clusters_A [Clusters_A$ Clustered_df_A == "1" ,])
# The centroid of cluster 1 is represented by a vector of mean values for each column of the data as a result.
```

```

Clustered_df_B <-cutree (hc_ward_cereal_b, k=5)
Clusters_B <-as.data.frame(cbind(cereal_b, Clustered_df_B))
Clust_2 <- colMeans (Clusters_B [Clusters_B$ Clustered_df_B == "1" ,])
# The centroid of cluster 2 is represented by a vector of mean values for each column of the data as a result.

```

```

Centroid <-rbind(Clust_1, Clust_2)
Centroid

```

```

##          calories    protein      fat    sodium      fiber    carbo
## Clust_1 -2.201871  1.3817478 -0.3310734  0.1727901  3.64131237 -2.0718749
## Clust_2  0.149818 -0.2449462  0.2483051 -0.2702057 -0.02091106 -0.7977876
##          sugars    potass    vitamins    weight      cups    rating
## Clust_1 -0.7894824  2.9837813 -0.1818422 -0.2008324 -1.845255  2.2426479
## Clust_2  1.0648712  0.1796942 -0.1818422  0.3369228 -0.303848 -0.5618826
##          Clustered_df_A
## Clust_1                1
## Clust_2                1

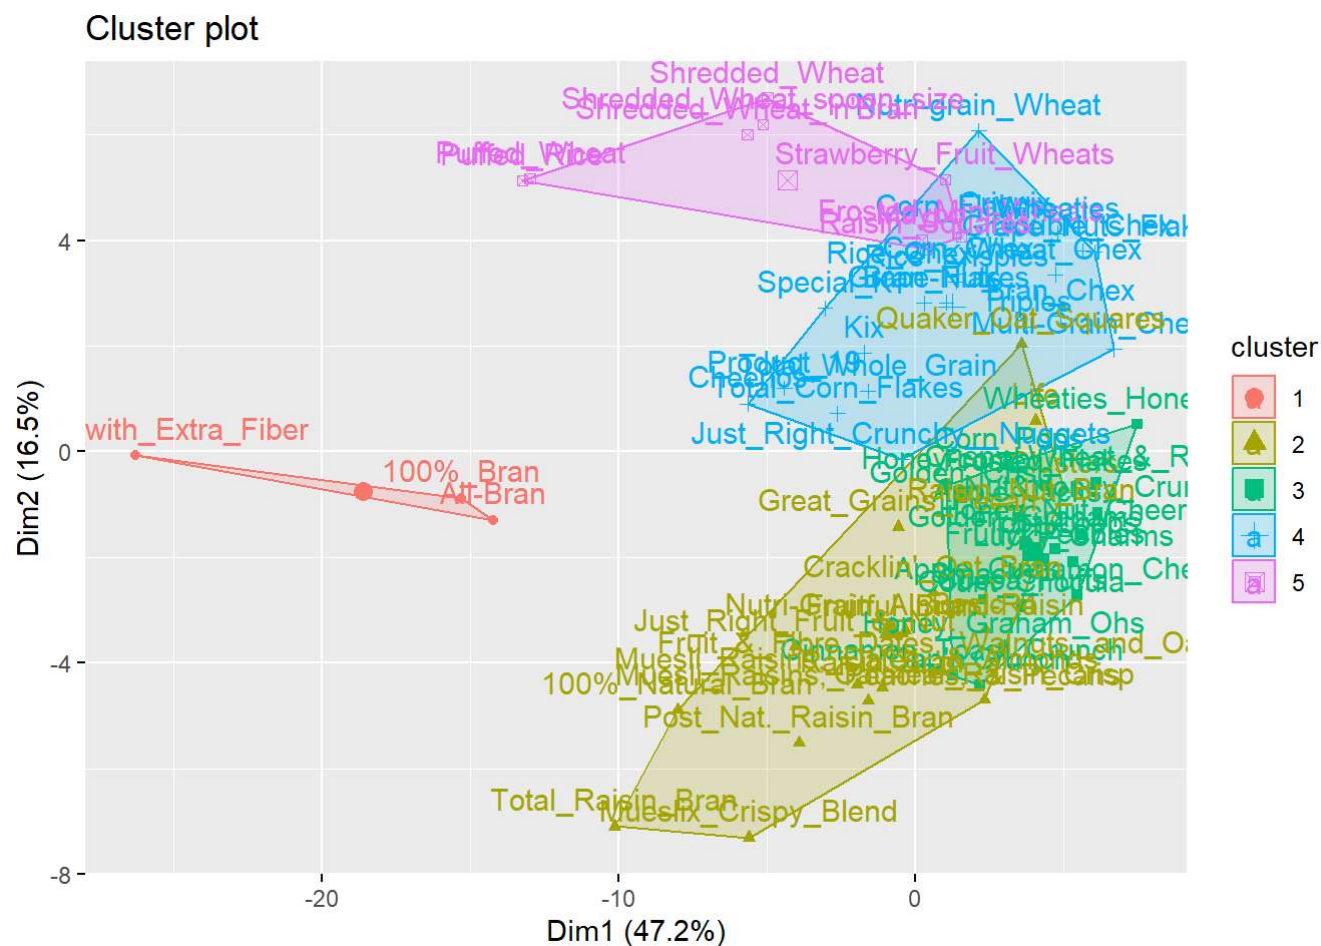
```

Question 3 (part C): Assess how consistent the cluster assignments are compared to the assignments based on all the data.

After reviewing the centroid, it shows that cluster 1 is high in protein, fiber, and potassium. It means that the cereals in cluster 1 are healthier than cluster 2. It can also be supported by looking at calories, fat, carbs, and sugar levels which are higher in cluster 2 as compared to cluster 1. Thus cereals in cluster 1 are healthier.

Q4: The elementary public schools would like to choose a set of cereals to include in their daily cafeterias. Every day a different cereal is offered, but all cereals should support a healthy diet. For this goal, you are requested to find a cluster of “healthy cereals.” Should the data be normalized? If not, how should they be used in the cluster analysis?


```
#Visualizing the clusters in Scatter plot
fviz_cluster(list(data=distance_table, cluster = c_d6))
```



```
Healthy_cereal<- cbind(c_d2,c_d6)
mean(Healthy_cereal[Healthy_cereal$c_d6==1,"rating"])
```

```
## [1] 73.84446
```

```
mean(Healthy_cereal[Healthy_cereal$c_d6==2,"rating"])
```

```
## [1] 38.37137
```

```
mean(Healthy_cereal[Healthy_cereal$c_d6==3,"rating"])
```

```
## [1] 28.66112
```

```
mean(Healthy_cereal[Healthy_cereal$c_d6==4,"rating"])
```

```
## [1] 46.17608
```



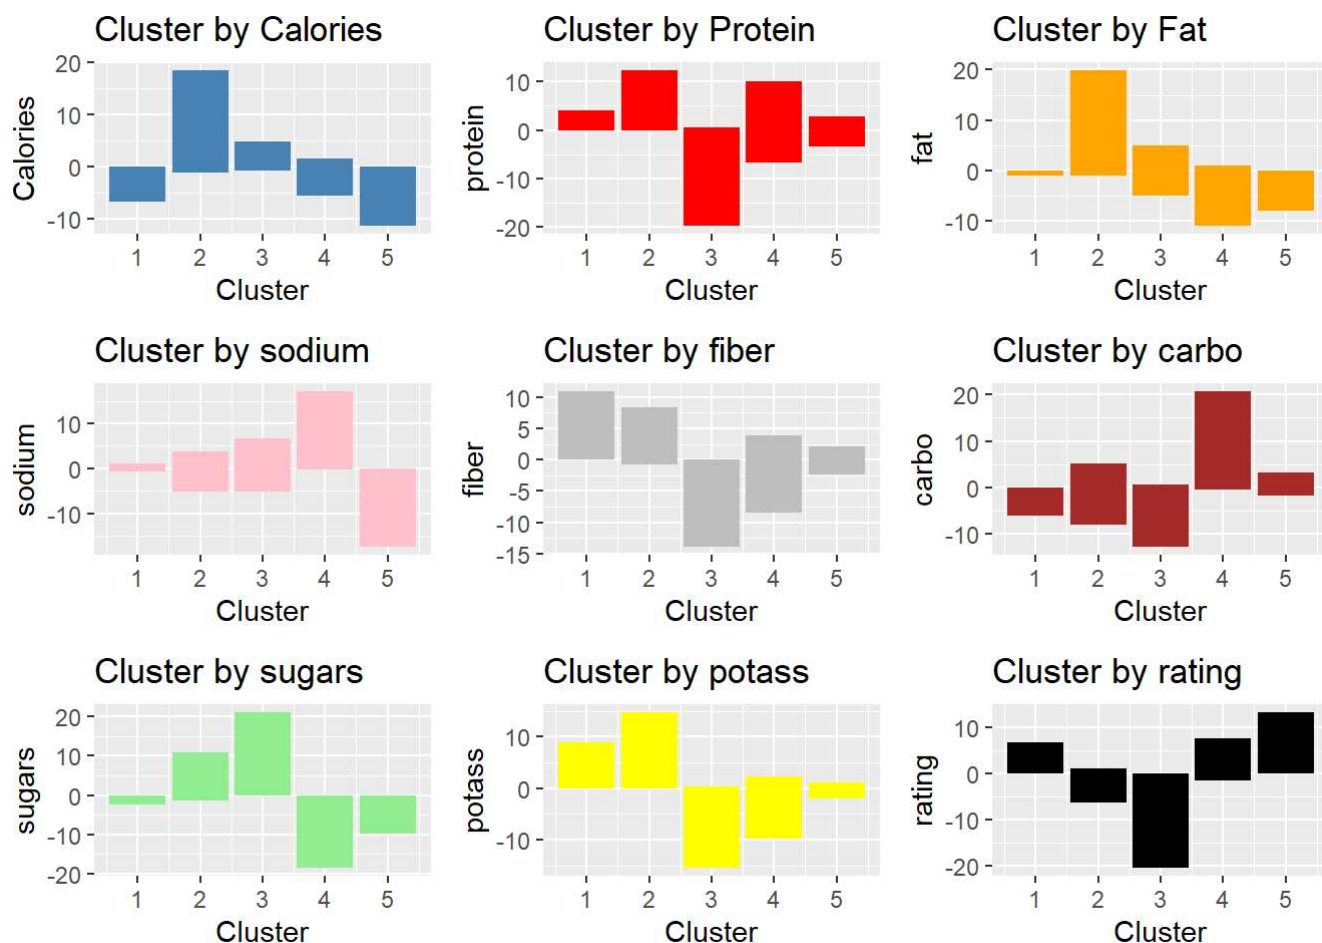
```
mean(Healthy_cereal[Healthy_cereal$c_d6==5,"rating"])
```

```
## [1] 63.0184
```

#It is evident that Cluster1 has the highest rating (73.84446), so we will select it as a nutritious cereal.

#lets also visualize the results by plotting a bar chart

```
calories <- ggplot(Clustered_df, aes(x = c_d6, y = calories)) +  
  geom_bar(stat = "identity", fill = "steelblue") +  
  labs(x = "Cluster", y = "Calories") +  
  ggtitle("Cluster by Calories")  
  
protein <- ggplot(Clustered_df, aes(x = c_d6, y = protein)) +  
  geom_bar(stat = "identity", fill = "red") +  
  labs(x = "Cluster", y = "protein") +  
  ggtitle("Cluster by Protein")  
  
fat <- ggplot(Clustered_df, aes(x = c_d6, y = fat)) +  
  geom_bar(stat = "identity", fill = "orange") +  
  labs(x = "Cluster", y = "fat") +  
  ggtitle("Cluster by Fat")  
  
sodium <- ggplot(Clustered_df, aes(x = c_d6, y = sodium)) +  
  geom_bar(stat = "identity", fill = "pink") +  
  labs(x = "Cluster", y = "sodium") +  
  ggtitle("Cluster by sodium")  
  
fiber <- ggplot(Clustered_df, aes(x = c_d6, y = fiber)) +  
  geom_bar(stat = "identity", fill = "gray") +  
  labs(x = "Cluster", y = "fiber") +  
  ggtitle("Cluster by fiber")  
  
carbo <- ggplot(Clustered_df, aes(x = c_d6,, y = carbo)) +  
  geom_bar(stat = "identity", fill = "brown") +  
  labs(x = "Cluster", y = "carbo") +  
  ggtitle("Cluster by carbo")  
  
sugars <- ggplot(Clustered_df, aes(x = c_d6,, y = sugars)) +  
  geom_bar(stat = "identity", fill = "lightgreen") +  
  labs(x = "Cluster", y = "sugars") +  
  ggtitle("Cluster by sugars")  
  
potass <- ggplot(Clustered_df, aes(x = c_d6,, y = potass)) +  
  geom_bar(stat = "identity", fill = "yellow") +  
  labs(x = "Cluster", y = "potass") +  
  ggtitle("Cluster by potass")  
  
rating <- ggplot(Clustered_df, aes(x = c_d6,, y = rating)) +  
  geom_bar(stat = "identity", fill = "black") +  
  labs(x = "Cluster", y = "rating") +  
  ggtitle("Cluster by rating")  
plot_grid(calories, protein, fat, sodium, fiber, carbo, sugars, potass, rating)
```



#Here we can see that cluster 1 still has the best results. It is low in calories, sugar, and fat. It has higher content of fiber, potassium and protein. Thus we can conclude that cluster 1 can be a set of cereals to include in their daily cafeterias.